

Project 18: Cycle Detection in a Graph

Sarvesh Soni
B.Math, 3rd Year

October 17, 2025

1 Problem Statement and Motivation

For this project, the main task was to implement an algorithm for detecting cycles in a graph. I learned that this is an important problem because cycles in a dependency graph can represent a "deadlock," where a set of processes are all waiting on each other and can't proceed. A simple example is in software package management, where if "Package A needs B, B needs C, and C needs A," it would make it impossible to install any of them. My task was to write a program that could find these situations.

2 Approach and Method Used

I decided to use a Depth-First Search (DFS) based algorithm. It felt like the most natural way to explore a graph path by path and check for "back edges," which are the main cause of cycles in a directed graph.

My implementation works by keeping track of the nodes in the current path of the traversal. For this, I used two boolean arrays:

- `visited[]`: This is just to keep track of all nodes I have ever visited. It helps to avoid doing redundant work on parts of the graph I have already checked and confirmed to be cycle-free.
- `recStack[]`: This is the important one. It keeps track of only the nodes that are in the current chain of recursive calls. If the DFS process ever lands on a node that is already in the `recStack`, it means I have followed a path back to an ancestor, which confirms a cycle.

The project specification also asked to compare with other methods. For an undirected graph, a much faster method is using a Disjoint Set (Union-Find) data structure. You can iterate through all edges and if you find an edge `'(u,v)'` where both `'u'` and `'v'` are already in the same set, you have found a cycle. This is very efficient but doesn't work for directed graphs, so I stuck with DFS.

3 Complexity Analysis

My analysis of the algorithm's complexity is as follows:

- **Time Complexity:** $O(V + E)$. This is the standard complexity for DFS. In the worst-case, my algorithm has to visit every single vertex (V) and cross every single edge (E) one time.

- Space Complexity: $O(V)$. The space is needed for the visited and recStack arrays. Both depend on the number of vertices. Also, the recursion call stack can go as deep as V in the worst case (like a single long chain graph).

4 Experiments and Datasets

To test my implementation on a real-world problem, I used the Bitcoin-Alpha trust network dataset from the Stanford Network Analysis Project (SNAP).

- Source: SNAP Website.
- Description: The dataset represents trust relationships. An edge from node A to B means user A trusts user B.
- Size: It has 3,783 nodes (users) and 24,186 edges (relationships).

A challenge with this dataset was that the node IDs are not small or sequential (like 0, 1, 2...). They go up to over 7000. So, my program first reads the whole file just to find the highest node ID. This way, I could create my graph structure with the exact size needed, instead of hardcoding a large number and wasting memory. Then, it reads the file a second time to actually add the edges. I treated each row as a directed edge from the 'SOURCE' column to the 'TARGET' column.

5 Results and Discussion

After loading the 24,186 edges from the dataset into my graph structure, I ran the cycle detection algorithm. The program produced a clear result:

RESULT: A cycle was detected in the graph.

The program also printed the first cycle it found to provide proof:

```
--- Cycle Found! ---
Cycle Path: 1 -> 160 -> 1
-----
```

This was an interesting finding. It means there are circular trust structures in the network (User 1 trusts User 160, and User 160 trusts User 1 back). This demonstrates that the algorithm works correctly on a non-trivial, real-world graph. The program ran very quickly, which aligns with the efficient $O(V+E)$ time complexity.

6 Challenges Faced and Lessons Learned

One of the first challenges was simply setting up the file reading logic in C++. It took some trial and error to correctly parse the CSV and handle potential errors, like improperly formatted lines.

The biggest lesson for me was seeing a classic textbook algorithm like DFS be applied to a real, messy dataset. It really solidified my understanding of how graphs can model real-world systems. Specifically, it drove home the difference between a general visited array and the recursionStack needed for tracking the current path. Getting this right was key. This project was a great experience in moving from theory to a practical, working program.